

Project requirements documentation

I want to build a project: symptoms and age-based medicine recommendation website—frontend: Next.js 14+ with tailwindcss and JavaScript - backend: FastAPI. I need full documentation and end to end requirements

Nice — great project. Below I'll give you a complete, research-grade **end-to-end requirements & documentation** package for a **symptoms + age-based medicine recommendation website** (Frontend: **Next.js 14+ + Tailwind + JavaScript**; Backend: **FastAPI**). It covers scope, safety & regulatory constraints, data sources, detailed architecture, API contracts, DB schema, ML / decision logic options, UX & legal text, security & privacy, testing, deployment, and an MVP checklist you can build from.

I cite authoritative sources for the most important legal / safety / data decisions. If you want, I can immediately convert any section into a runnable repo scaffold or detailed code files next.

1 — Project summary & scope (what this system will and will not do)

Purpose (high level)

Provide age- and symptom-based *recommendations* for over-the-counter (OTC) medicines and self-care guidance — and flag serious conditions that require urgent medical attention. The system should **NOT** replace professional medical diagnosis**. ** It must include clear disclaimers, triage rules (red flags), and encourage seeing a clinician where appropriate.

Core features (MVP)

- Symptom intake form (age, weight optionally, gender, pregnancy status, allergies, current meds).
- Rule-based triage (red/amber/green).
- Age-aware medicine suggestions for simple, minor complaints (e.g., paracetamol for fever in adult vs child dosing *if allowed by local regulation*).

- Drug interaction & contraindication checks (using authoritative drug vocabularies/APIs).
- Clear safety / warning UI and hard stops for "seek urgent care".
- Audit log of recommendations, versioned clinical rules, and consent capture.

Out of scope (initially)

- Prescription medication issuance.
- Definitive diagnoses.
- Replacing clinician workflows.
- Automated dispensing / medication delivery (unless explicitly required later and vetted legally).

Safety note: This is a high-risk domain. Follow clinical decision support (CDS) guidance and privacy regulations. See the FDA CDS guidance and HIPAA/GDPR sections below.

2 — Regulatory & safety constraints (must-read before design)

- 1. Clinical Decision Support regulation (US / FDA):** determine if your recommendation engine falls under FDA definition of a medical device / CDS — design to meet “non-device CDS” criteria or prepare for regulatory requirements.
- 2. Protected health information (HIPAA) if US users or US covered entities involved:** If you store, transmit, or process identifiable health information and you are a covered entity or business associate, HIPAA technical & administrative safeguards apply (encryption at rest/in transit, access controls, breach notifications, auditing).
- 3. EU GDPR — special category data (health):** health data is a “special category” and requires lawful basis + additional safeguards (explicit consent, data minimization, DPIA). If you operate in/serve EU users, GDPR obligations are mandatory.
- 4. Use authoritative drug data & licensing:** For dosing, interactions, contraindications use RxNorm, DrugBank, DailyMed/openFDA or other region-appropriate sources. Watch commercial license terms (DrugBank often needs subscription for clinical API).

5. **User safety & UX:** implement conservative triage (prefer false positives for escalation), always show “see a clinician” for uncertainty, and keep logs/versioning of rules for auditability. See WHO guidance on self-care and published work on chatbot security in health.
-

3 — Data sources & licensing (recommended)

Primary drug vocabularies / clinical data

- **RxNorm** (US normalized drug names, APIs): use for mapping drug names, active ingredients, dose forms.
- **DrugBank Clinical API** (drug interactions, contraindications — commercial terms).
- **openFDA / DailyMed** (US drug labels, adverse events, dosing sections).

Symptom → condition / triage knowledge

- Curated clinical rule sets you author or adapt from public guidance (WHO self-care, NHS symptom checkers when license allows). Avoid using conversational LLMs for final rule decisions without clinician review.

Other

- ICD-10 or SNOMED CT for coding conditions (be mindful of SNOMED license).
- Local medication formularies and pediatric dosing tables (region dependent).

Licensing note: validate license & commercial use (DrugBank may require paid license for clinical module; RxNorm is public). Always record data source and version for each recommendation.

4 — High-level architecture

Diagram (conceptual):

lua

```
[Next.js Frontend] <--> [FastAPI Backend] <--> [DB (Postgres + Auditing)]
|
+--> Drug data cache (Redis) / RxNorm & DrugBank API calls
+--> Rules engine (Drools / simple Python rule module)
```

+--> *ML service (optional) (separate microservice)*

+--> *Logging / Audit & Monitoring (ELK / Prometheus)*

Components detail

- **Frontend (Next.js 14+, JS, Tailwind):** symptom pages, age/dose forms, consent modal, results UI, warnings, account pages.
- **Backend (FastAPI):** REST + OpenAPI, auth (OAuth2 / JWT), business logic, rules engine, connectors to drug APIs.
- **DB:** PostgreSQL (encrypted at rest); separate table for anonymized analytics; store only minimal identifiable PHI if required.
- **Cache:** Redis for drug lookups / interaction checks.
- **Rules engine:** deterministic rule set to do triage and age/dose logic. Keep rules versioned in DB or as configuration files.
- **ML (optional):** a classifier (symptoms → probable condition) used only to *suggest* and always surfaced with confidence and rule overrides; final recommendations must be rule-checked.
- **Audit & Monitoring:** immutable logs for every recommendation, who requested it, rule version, data sources used, timestamp. Use WORM logs or append-only DB schema where regulation requires.

5 — Security, privacy & compliance (practical checklist)

Technical safeguards

- TLS 1.2+ everywhere; HSTS.
- Encryption at rest (DB encryption + field-level encryption for PHI).
- Access control: RBAC and least privilege; admin access via MFA.
- JWT for session tokens; rotate keys.
- Rate limiting + WAF (Cloudflare / AWS WAF).
- Secrets management (Vault / AWS Secrets Manager).
- Audit logging & immutable event store for recommendations.

Privacy

- Minimal data collection; anonymize analytics.
- Obtain **explicit consent** for processing health data. Record consent (who, when, for what).
- Data retention policy & deletion flow (user-initiated deletion).

- Data Processing Agreement (DPA) with cloud providers if you process EU/EAA data.

Organizational

- Appoint Data Protection Officer (GDPR) if required.
- Conduct a DPIA for health data processing.
- Incident response & breach notification procedures (HIPAA: notify within 60 days; follow local law).

(References: HIPAA & GDPR summaries).

6 — Clinical safety design: rules, triage, disclaimers

Triage strategy (recommended conservative approach)

- **Red flags** (immediate referral to emergency): chest pain, severe shortness of breath, loss of consciousness, severe bleeding, suspected stroke signs, severe head injury, very high fever in infants.
- **Amber**: moderate symptoms, vulnerable populations (very young, elderly, pregnant) → recommend clinician contact.
- **Green**: self-care / OTC options with dosing instructions and warnings.

Rule engine design

- Rules expressed declaratively (YAML or JSON) with versioning. Example rule fields: `id`, `name`, `conditions` (symptoms, duration, vitals if provided), `age_range`, `pregnancy_flag`, `recommended_actions`, `priority`, `source` (e.g., WHO, local guideline), `created_by`, `version`.
- Each recommendation includes `evidence_source` and `confidence_score` and lists contraindications & interactions evaluated.

Medication checks

- Confirm: age, weight (if pediatric dosing), pregnancy status, allergies, current meds.
- Query drug interaction APIs and contraindication lists. If interaction exists, **do not** recommend and show alternatives. Use a conservative block for serious interactions.

Disclaimer + informed consent

- Prominent banner: "This tool provides information only, not a diagnosis. For emergencies call local emergency services. Always follow your clinician's advice."

- Before collecting health data, show consent modal explaining storage, retention, and purpose.

(See FDA CDS guidance and WHO self-care resources).

7 — API specification (core endpoints)

All endpoints return standard problem+json for errors. Use OpenAPI (FastAPI auto-doc).

Auth

- `POST /auth/login` → returns JWT (access token, refresh token).
- `POST /auth/register` → sign up (capture consent).

Symptom intake & triage

- `POST /api/v1/symptoms/assess`
 - Body:

json

```
{
  "user_id": "uuid or null",
  "age_years": 35,
  "age_months": null,
  "weight_kg": 70,
  "sex": "female",
  "pregnant": false,
  "symptoms": ["fever", "headache"],
  "symptom_onset_days": 2,
  "current_medications": ["ibuprofen"],
  "allergies": ["penicillin"],
  "location_country": "IN"
}
```

- Response (abridged):

json

```
{
  "triage": "amber",
  "reasons": ["age > 65", "symptom severity: moderate"],
}
```

```

"recommendations": [
  {
    "id": "rec_001",
    "title": "Paracetamol 500 mg oral",
    "dose": "500 mg every 4–6 hours; max 3 g/day for adults",
    "age_restrictions": ">=12 years",
    "contraindications": ["severe liver disease"],
    "drug_interactions": [],
    "evidence_sources": ["RxNorm vX.Y", "Local formulary v1"]
  }
],
"rule_version": "2025-10-01-v3",
"audit_id": "uuid"
}

```

Drug details

- `GET /api/v1/drug/{rxnorm_id}` → returns drug metadata, interactions, pediatric dosing.

Audit

- `GET /api/v1/audit/{audit_id}` → returns logged inputs, rules used, sources.

Admin

- `POST /api/v1/rules` → upload/patch rule sets (admin only).
- `GET /api/v1/rules?version=latest` → view applied rules.

8 — Database schema (core tables)

Simplified (Postgres):

- `users` (id, email, hashed_password, consent_recorded_at, country, created_at)
- `assessments` (id, user_id nullable, input_json, triage_result, recommendations_json, rule_version, audit_hash, created_at)
- `rules` (id, name, yaml/json, version, author, created_at)
- `drugs` (rxnorm_id, name, active_ingredient, atc_code, last_synced, source)
- `drug_interactions` (id, drug_a, drug_b, severity, source)
- `audit_logs` (id, assessment_id, action, actor, timestamp, signature)

Store encrypted `input_json` when it contains identifiers/PHI.

9 — Rules authoring & governance

- Use a clinician-reviewed process: author → peer review (2 clinicians) → QA test suite (automated scenario tests) → publish with semantic versioning.
 - Maintain a **rules rollback** and **changelog**.
 - Each recommendation must link to an evidence source (URL + snapshot of the relevant text) and the rule unit tests that validate it.
-

10 — ML vs Rule-based: recommended hybrid

- **Rules-first (mandatory)**: Triage & medication checks must be deterministic rules.
- **ML (optional assistive layer)**: symptom → probable condition ranking (for UX to suggest likely causes). Never auto-apply ML output without rule-based validation. Log model version and confidence. Implement guardrail: if model confidence < threshold, show only rule outputs or escalate.

Model training data: clinical symptom datasets are sensitive; use de-identified public datasets or simulated data; never train on PHI.

11 — Frontend structure & UX patterns (Next.js 14+ + Tailwind)

Pages / flows

- `/` Landing + simple search symptom box.
- `/assess` Symptom intake wizard (stepper): basic demographics → symptoms → current meds/allergies → review & consent → results.
- `/results/[audit_id]` Result page with clear triage level, recommendations, what to do next.
- `/account` (optional): saved assessments, share/export, delete data.
- Admin UI (separate app or protected path): rule management, audit viewer.

Key UX elements

- Large, color-coded triage banners (red/amber/green) with accessible text.
- “Why this recommendation?” modal listing evidence & check results (interactions, contraindications).
- Print / PDF export of assessment with audit id for clinician.

- Explicit “Not a diagnosis” header in results and consent screen.

Accessibility: WCAG AA, large font, ARIA labels for screen readers, keyboard navigation.

12 — Example FastAPI snippet (skeleton)

python

```
# main.py (abridged)
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel
import uuid, datetime

app = FastAPI(title="SymptomMed API")

class AssessmentIn(BaseModel):
    age_years: int
    symptoms: list[str]
    current_medications: list[str] = []
    allergies: list[str] = []
    pregnant: bool = False
    # ... add fields

@app.post("/api/v1/symptoms/assess")
async def assess(payload: AssessmentIn):
    # 1. basic validation
    if payload.age_years < 0:
        raise HTTPException(400, "Invalid age")
    # 2. run rule engine (placeholder)
    triage = "green" # real logic to replace
    recommendations = [] # call drug rule checks
    audit_id = str(uuid.uuid4())
    # store assessment (save input + result) - omitted
    return {
        "triage": triage,
        "recommendations": recommendations,
        "rule_version": "2025-10-01-v1",
        "audit_id": audit_id
    }
```

I can expand this into a full FastAPI project structure if you want.

13 — Testing & QA

Automated tests

- Unit tests for rule conditions.
- Integration tests: full assessment flows, API → DB → drug API mocks.
- End-to-end tests (Playwright/Cypress) for important user journeys (consent, assessment, result).
- Security tests (SAST / dependency scanning).
- Penetration test before production if PHI involved.

Clinical validation

- Create a test suite of clinical vignettes (with clinician input) that assert correct triage & meds. Maintain coverage > 95% for core rules.
-

14 — Monitoring, logging & observability

- **Metrics:** assessments/day, triage distribution, API latency, error rates.
 - **Logging:** structured logs with correlation IDs; store audit entries separately.
 - **Alerting:** high error rate, data exfiltration signs, sudden change in triage distribution.
 - **Model / rule drift:** track frequency of rule overrides and clinician feedback.
-

15 — Deployment & infra (recommended stack)

- Containerize (Docker).
- Deploy via managed K8s (EKS/GKE/AKS) or App Platform. Use horizontal autoscaling.
- PostgreSQL (managed RDS / Cloud SQL) with encryption at rest + daily backups.
- Redis for caching.
- Use a secrets manager and IAM roles for service auth.
- CI/CD: GitHub Actions / GitLab CI — run tests, security scans, then deploy to staging & production with approvals.
- Use feature flags for new rules (LaunchDarkly / Unleash) to roll out safely.

16 — Logging & audit retention policy (example)

- Retain assessment logs for **minimum** 7 years if required by local law (medical records retention varies). Provide user-controlled deletion for non-required jurisdictions. Keep an immutable audit log for compliance, masked where necessary.
-

17 — Team & roles (who you'll need)

- Product owner / clinical lead (licensed clinician).
 - Backend engineer (FastAPI, Postgres, DevOps).
 - Frontend engineer (Next.js 14+, Tailwind).
 - Data engineer (drug APIs, ETL).
 - QA engineer + clinical QA.
 - Security & privacy officer / legal counsel (for HIPAA/GDPR).
 - DevOps / SRE for infra and monitoring.
-

18 — MVP checklist (prioritized)

Phase 1 (MVP)

- Symptom intake + consent flow.
- Rule-based triage & one common condition (e.g., fever/cold) recommendations (adult only).
- RxNorm integration for drug metadata.
- Audit logging & rule versioning.
- Visible disclaimers & emergency red flags.
- Unit & integration tests for rules.

Phase 2

- Pediatric dosing (weight-based).
- Drug interaction checks (DrugBank or openFDA).
- Auth / accounts, export PDF.
- Clinical review pipeline & rule editor.

Phase 3

- Multi-region regulatory compliance, ML assistant (probable diagnostics), BI dashboards, clinician feedback loop.
-

19 — Example risk register (top risks + mitigations)

1. **Incorrect medication suggestion** — mitigate: clinician review, conservative rules, block serious interactions.
 2. **Data breach of PHI** — mitigate: encryption, access controls, incident response plan.
 3. **Regulatory enforcement (FDA / local)** — mitigate: legal review, avoid features that classify as regulated medical device unless you plan to comply.
-

20 — Sources & further reading (most-load bearing)

- FDA — Clinical Decision Support Software guidance.
 - HHS — HIPAA Privacy & Security Rule summaries.
 - GDPR — Article 9 (special categories / health data).
 - RxNorm / RxNav APIs — drug normalization & APIs.
 - DrugBank Clinical API (interactions data; commercial).
 - openFDA / DailyMed — public drug label and event data.
 - WHO — self-care for health & minor ailments guidance.
 - Research on AI chatbot security in healthcare.
-

21 — Next steps (practical choices for you right now)

Tell me which of these you want immediately and I'll produce it in this session (I will deliver now, not later):

1. Full **FastAPI** project scaffold (Dockerfile, endpoints, rule engine stub, tests).
2. Full **Next.js 14+** scaffold (pages, Tailwind config, assessment UI components, example integration).
3. A **rules pack** for 10 common symptoms (YAML) + unit tests and clinical vignette suite.
4. A **detailed security & privacy checklist** customized to the countries you'll serve (e.g., India + EU + US).

5. A sample contract / consent wording and privacy notice suitable for HIPAA/GDPR audiences.

Pick one (or more) and I'll generate the requested artifact immediately — e.g., a ready-to-run repo scaffold or rules YAML + tests.